

Efficient Pairing Functions—and Why You Should Care

Arnold L. Rosenberg*

*Dept. of Computer Science, Univ. of Massachusetts
Amherst, MA 01003, USA*

ABSTRACT

This paper provides a short guided tour through the world of *pairing functions*—bijections between $\mathbb{N} \times \mathbb{N}$ and \mathbb{N} —as models for computational “situations.” After a short discussion of the computationally simplest pairing functions—the Cauchy-Cantor “diagonal” polynomials—we describe in detail two specific computational situations where pairing functions give an unexpectedly simple handle on apparently complex systems-related problems. (1) We discuss the use of pairing functions as storage mappings for rectangular arrays/tables that are *extendible*, in the sense that the programmer may expand and shrink them dynamically. (2) We discuss the use of pairing functions as the basis for a mechanism for instilling *accountability* into Web-computing projects, by linking remote “volunteers” to the tasks they compute.

Keywords: Pairing function, storing extendible arrays/tables, accountability in web computing

1. Introduction

1.1. Motivation

Entia non sunt multiplicanda praeter necessitatem.

Occam’s Razor

This famous admonition by William of Occam (14th cent.) to strive for simplicity is worth heeding when seeking mathematical models of computational phenomena. In this spirit, we describe examples of the use of *pairing functions* (PFs, for short)—bijections between $\mathbb{N} \times \mathbb{N}$ and \mathbb{N} , where \mathbb{N} is the set of positive integers—as the basis for such models. The utility of PFs in myriad disparate situations resides in their allowing one to slip gracefully between one- and two-dimensional worldviews—and, by iteration, among worldviews of arbitrary finite dimensionalities. Our emphasis is on the varied structures that PFs can enjoy, which allow one to use PFs to model a broad range of quite disparate computational situations, including their associated natural measures of efficiency. After due obeisance to the importance of PFs in foundational studies (in the next subsection), we turn in Section 2 to a short discussion of the computationally simplest PFs, the Cauchy-Cantor polynomials, whose charming mathematical structure—see (2.1)—begs one to ask if other such polynomials exist. This question remains unresolved, but there

*rsnrg@cs.umass.edu

has been nontrivial progress toward a (thus far) negative answer. In the remainder of the paper, we focus on PFs in more “modern” computational situations, discussing just two situations that will hopefully convince the reader of the utility of PFs as a modeling device. (1) We discuss the use of PFs as storage mappings for rectangular arrays/tables that are *extendible*, in that the programmer can expand and shrink them dynamically (Section 3); our focus is on storage mappings that store arrays/tables “compactly.” (2) We discuss the use of PFs as the basis for a mechanism for instilling *accountability* into Web-computing projects, by linking “volunteers” to the tasks they compute (Section 4); our focus is on PFs that map each “row” of $\mathbb{N} \times \mathbb{N}$ to an arithmetic progression. During our short tour, we shall encounter several intriguing and challenging problems that remain open, despite the apparent simplicity of PFs.

1.2. Background

PFs have played a major role in a variety of “classical” studies. They played a pivotal role in Cantor’s seminal study of infinities [1], supplying a rigorous formal basis for asserting the counterintuitive “equinumerousness” of the integers and the rationals. It took revolutionary thinkers such as Gödel and Turing to recognize that the correspondences embodied by PFs can be viewed as *encodings*, or *translations*, of ordered pairs (and, thence, of arbitrary finite tuples or strings) as integers. This insight allowed Gödel and Turing to build on the existence of eminently computable—indeed, easily computed—PFs in their famous studies of, respectively, logical systems [5] and algorithmic systems [17]. The uses we propose for PFs, while certainly less profound than these two, also build on the insight that PFs can be used as encoding mechanisms, specifically allowing one to slip gracefully among worlds of strings, integers, and tuples of integers.

Throughout the paper, we illustrate selected values from selected PFs using the following convention. We illustrate a PF $\mathcal{F} : \mathbb{N} \times \mathbb{N} \leftrightarrow \mathbb{N}$ via a two-dimensional array whose entries are the values of \mathcal{F} as described in Fig. 1.

$\mathcal{F}(1, 1)$	$\mathcal{F}(1, 2)$	$\mathcal{F}(1, 3)$	$\mathcal{F}(1, 4)$	$\mathcal{F}(1, 5)$	\dots
$\mathcal{F}(2, 1)$	$\mathcal{F}(2, 2)$	$\mathcal{F}(2, 3)$	$\mathcal{F}(2, 4)$	$\mathcal{F}(2, 5)$	\dots
$\mathcal{F}(3, 1)$	$\mathcal{F}(3, 2)$	$\mathcal{F}(3, 3)$	$\mathcal{F}(3, 4)$	$\mathcal{F}(3, 5)$	\dots
$\mathcal{F}(4, 1)$	$\mathcal{F}(4, 2)$	$\mathcal{F}(4, 3)$	$\mathcal{F}(4, 4)$	$\mathcal{F}(4, 5)$	\dots
$\mathcal{F}(5, 1)$	$\mathcal{F}(5, 2)$	$\mathcal{F}(5, 3)$	$\mathcal{F}(5, 4)$	$\mathcal{F}(5, 5)$	\dots
\vdots	\vdots	\vdots	\vdots	\vdots	\ddots

Figure 1: *Our generic template for sampling from a PF.*

2. The Prettiest Pairing Function(s)

It has been known for almost two centuries that there exist bijections between $\mathbb{N} \times \mathbb{N}$ and \mathbb{N} . Indeed, it has been known for at least 125 years that there exist

1	3	6	10	15	21	28	36	...
2	5	9	14	20	27	35	44	...
4	8	13	19	26	34	43	53	...
7	12	18	25	33	42	52	63	...
11	17	24	32	41	51	62	74	...
16	23	31	40	50	61	73	86	...
22	30	39	49	60	72	85	99	...
29	38	48	59	71	84	98	113	...
⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮

Figure 2: The diagonal PF \mathcal{D} . The shell $x + y = 6$ is highlighted.

such bijections that are *polynomials*. Cauchy [2] pictorially describes, and Cantor [1] symbolically specifies, the *Diagonal* PF

$$\mathcal{D}(x, y) = \binom{x + y - 1}{2} + y \quad (2.1)$$

(which, of course, has a twin obtained by exchanging x and y); see Fig. 2. A simple double induction based on the fact that \mathcal{D} maps integers in an “upward direction” along the “diagonal shells,” $x + y = 2$, $x + y = 3$, $x + y = 4$, ..., proves that the function \mathcal{D} is, indeed, a bijection. A computationally more satisfying proof in [3] presents an explicit recipe for computing \mathcal{D} ’s inverse.

Is \mathcal{D} the only polynomial PF? It is so intriguing that there exists a PF that is a *polynomial* that the question of \mathcal{D} ’s uniqueness (as a *polynomial* PF) is irresistible! This question remains largely open, but there are a few nontrivial beginnings to an answer.

1. There is no quadratic polynomial PF other than \mathcal{D} (and its twin) [4].
2. The preceding assertion remains true if the “onto” condition for bijections is replaced by a “unit density” condition [7].
3. No cubic or quartic polynomial is a PF [8].
4. The development in [8] excludes large classes of higher-degree polynomials from being PFs—a simple example: a super-quadratic polynomial whose coefficients are all positive cannot be a PF.

The proofs of the preceding results combine geometric number theory with non-standard diophantine analysis. (Standard diophantine analysis studies the existence of *any* integer solution to a polynomial equation and the existence of *infinitely many* integer solutions; it does not address the existence of *a unique* integer solution.) For instance, the lead terms of any super-quadratic polynomial \mathcal{F} grow faster than the quadratic growth of the plane, hence must leave large gaps in their ranges. To prove

that such an \mathcal{F} cannot be a PF, one must show that \mathcal{F} 's lower-degree terms do not lead to large “troughs” that capture all of the integers that its lead terms miss. The cases that must be eliminated grow quickly in number with the degree of \mathcal{F} —cf. [8]—yet in principle provide no impenetrable barrier to the ultimate resolution of the question.

3. Pairing Functions and Array/Table Storage

It has been recognized since the 1950s that many classes of computations, ranging in origin from linear-algebraic scientific applications to (relational) databases, benefit from the ability to reshape multidimensional arrays and tables dynamically. While several programming languages allow a user to specify at least some types of reshapings—say, the addition and/or deletion of rows and/or columns in two dimensions—the language processors I am aware of implement the capability quite naively, by completely remapping an array/table with each reshaping. This is, of course, very wasteful of time, since one does $\Omega(n^2)$ work to accommodate $O(n)$ changes.

In the mid-1970s, I began to study the use of PFs as storage-mapping functions for two-dimensional rectangular arrays/tables, since the mappings so specified would allow one to add and/or delete rows and/or columns dynamically, without ever remapping array/table positions that are unaffected by the reshaping. (Extending this work to higher dimensionalities is immediate.) This section surveys some of the results I obtained.

3.1. A Methodology for Constructing PFs

Rather than trying to select a single one-form-fits-all storage mapping for dynamically extendible arrays, I decided, in [11], to explore the space of possible PFs, with an eye to seeking ones that optimized various criteria. I used the following procedure to systematize the process of constructing PFs.

Procedure PF-Constructor(\mathcal{A})
/*Construct a PF \mathcal{A}^* /*

Step 1. Partition the set $\mathbb{N} \times \mathbb{N}$ of potential array/table positions into finite sets called *shells*. Order the shells linearly in some way: many natural shell-partitions carry a natural order.

Aside. Here are a few sample shell-partitions that played a role in our study of extendible arrays. For each relevant integer c , shell c comprises all pairs $\langle x, y \rangle$ such that:

$x + y = c$	the diagonal shells that define the PF \mathcal{D} of (2.1) and Fig. 2
$\max(x, y) = c$	square shells; cf. (3.3) and Fig. 3
$xy = c$	hyperbolic shells; cf. (3.4) and Fig. 4.

Step 2. Construct a PF from the shells as follows.

Step 2a. Enumerate the array positions shell by shell, honoring the ordering of the shells.

Step 2b. Enumerate each shell in some systematic way, say “by columns.”

This means enumerating the pairs $\langle x, y \rangle$ in the shell in increasing order of y and, for pairs having equal y values, in, say, decreasing order of x . (Increasing order of x works as well, of course.)

Theorem 3.1 *Any function $\mathcal{A} : \mathbb{N} \times \mathbb{N} \leftrightarrow \mathbb{N}$ that is designed via Procedure PF-Constructor is a valid APF.*

Proof. We sketch the easy proof. Step 1 of Procedure PF-Constructor constructs a partial order on $\mathbb{N} \times \mathbb{N}$, in which: (a) each set of incomparable elements—called a shell—is finite; (b) there is a linear order on the shells. Step 2 extends the partial order of Step 1 to a linear order, by honoring the linear order on the shells and imposing a linear order within each shell. The function \mathcal{A} constructed by the Procedure can thus be viewed as an enumeration of $\mathbb{N} \times \mathbb{N}$ —which means that \mathcal{A} is a PF. ■

Of course, Procedure PF-Constructor begs the question of how to compute the PF \mathcal{A} efficiently. This issue will be a major concern for the remainder of this section.

3.2. Pursuing Compact PFs

When one considers using a PF for mapping arrays/tables into storage, one notes immediately the poor resulting management of storage. For instance, the diagonal PF \mathcal{D} spreads the n^2 -position $n \times n$ array/table over $2n^2$ addresses: $\mathcal{D}(1, 1) = 1$ and $\mathcal{D}(n, n) = 2n^2$; even worse (percentage-wise), \mathcal{D} spreads the n -position $1 \times n$ array/table over $> \frac{1}{2}n^2$ addresses: $\mathcal{D}(1, 1) = 1$ and $\mathcal{D}(1, n) = \frac{1}{2}(n^2 + n)$. This loss of “compactness” is a more serious deficiency than is the loss of the bidirectional arithmetic progressions enjoyed by the standard row- or column-major indexings used by most compilers, since the waste of storage plagues one no matter how one intends to access the array/table. Therefore, in [11, 12], I actively sought PFs \mathcal{A} that were *compact*, as measured by small growth rates of their *spread functions*

$$\mathbf{S}_{\mathcal{A}}(n) \stackrel{\text{def}}{=} \max\{\mathcal{A}(x, y) \mid xy \leq n\}. \quad (3.1)$$

In other words, $\mathbf{S}_{\mathcal{A}}(n)$ is the largest address that the PF \mathcal{A} assigns to any position of an array/table that has n or fewer positions.

Here is what I discovered.

3.2.1. PFs that favor one fixed aspect ratio

Say that one moderates one’s demands on the PF \mathcal{A} by focusing only on its compactness when storing arrays/tables of a fixed aspect ratio $\langle a, b \rangle$, i.e., arrays/tables whose dimensions have the form $ak \times bk$ for some k . Then one can manage storage perfectly, in the sense that there exists a PF $\mathcal{A}_{a,b}$ such that

$$\mathbf{S}_{\mathcal{A}_{a,b}}(n) \stackrel{\text{def}}{=} \max\{\mathcal{A}_{a,b}(x, y) \mid [x \leq ak] \wedge [y \leq bk] \wedge [abk^2 \leq n]\} = n. \quad (3.2)$$

In other words, $\mathcal{A}_{a,b}$ maps every position (x, y) of an $ak \times bk$ array/table having n or fewer positions to an address $\leq n$. It is easy to construct $\mathcal{A}_{a,b}$ via the shells specified as follows.

1. Shell 1 comprises the positions of the $a \times b$ array, i.e., the set

$$\{\langle x, y \rangle \mid [x \leq a] \wedge [y \leq b]\}.$$

2. Shell $k + 1$ comprises the positions of the $a(k + 1) \times b(k + 1)$ array that are not elements of the $ak \times bk$ array, i.e., the set

$$\{\langle x, y \rangle \mid [ak < x \leq a(k + 1)] \vee [bk < y \leq b(k + 1)]\}.$$

While the preceding guarantee of compactness ignores the issue of ease of computation, there are at least some PFs that utilize storage perfectly in the sense of (3.2) and are quite easy to compute. One useful such PF favors *square* arrays/tables, i.e., those whose aspect ratio is given by $a = b = 1$. This “square-shell” PF, $\mathcal{A}_{1,1}$ is specified by the following explicit expression.

$$\begin{aligned} \mathcal{A}_{1,1}(x, y) &= m^2 + m + y - x + 1 \\ \text{where } m &\stackrel{\text{def}}{=} \max(x - 1, y - 1). \end{aligned} \quad (3.3)$$

Once having noted that $\mathcal{A}_{1,1}$ maps integers in a counterclockwise direction along the “square shells” $m = 0, m = 1, \dots$ (see Fig. 3), one verifies its bijectiveness via a simple double induction. (Of course, $\mathcal{A}_{1,1}$ has a twin that proceeds in a clockwise direction along the square shells.)

1	4	9	16	25	36	49	64	...
2	3	8	15	24	35	48	63	...
5	6	7	14	23	34	47	62	...
10	11	12	13	22	33	46	61	...
17	18	19	20	21	32	45	60	...
26	27	28	29	30	31	44	59	...
37	38	39	40	41	42	43	58	...
50	51	52	53	54	55	56	57	...
⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮

Figure 3: The square-shell PF $\mathcal{A}_{1,1}$. The shell $\max(x, y) = 5$ is highlighted.

3.2.2. PFs that favor finite sets of aspect ratios

We showed in [12] how to “dovetail” any set $\{\mathcal{A}_1, \mathcal{A}_2, \dots, \mathcal{A}_m\}$ of m PFs to arrive at a PF \mathcal{A} whose compactness is at worst m times that of the most compact of the PFs \mathcal{A}_i ; i.e., for all n ,

$$\mathbf{S}_{\mathcal{A}}(n) \leq m \cdot \min_i \mathbf{S}_{\mathcal{A}_i}(n).$$

The dovetailing is performed in two steps.

1. Alter each \mathcal{A}_k to be a bijection $\mathcal{A}_k^{(m)}$ between $\mathbb{N} \times \mathbb{N}$ and the congruence class $(k - 1) \bmod m$, i.e., the set of integers of the form $mx + k - 1$. Specifically, define the PF $\mathcal{A}_k^{(m)}$ via: $\mathcal{A}_k^{(m)}(x, y) = m \cdot \mathcal{A}_k(x, y) + k - 1$.
2. Define the PF \mathcal{A} via: $\mathcal{A}(x, y) = \min_k \{\mathcal{A}_k^{(m)}(x, y)\}$.

Thus, if one wants a PF to be compact on arrays of any fixed finite set of, say m , aspect ratios $\langle a_1, b_1 \rangle, \langle a_2, b_2 \rangle, \dots, \langle a_m, b_m \rangle$, then one can craft a PF $\mathcal{A}_{a_1, b_1; a_2, b_2; \dots; a_m, b_m}$ that maps every position (x, y) of an array/table which has one of these m aspect ratios, and which has n or fewer positions, to an address $\leq mn$.

3.2.3. A PF that minimizes worst-case spread

The preceding shape-based guarantees do not help much with applications such as relational databases, wherein one cannot limit *a priori* the potential shapes of one's tables. This fact led me to the question of how compact a PF could be on arrays/tables of arbitrary shapes. I showed in [12] that there exists a PF \mathcal{H} whose compactness is given by

$$\mathbf{S}_{\mathcal{H}}(n) = O(n \log n),$$

and no PF can beat this compactness (in the worst case) by more than a constant factor. The PF \mathcal{H} maps integers along the “hyperbolic shells” defined by $xy = 1, xy = 2, xy = 3, \dots$. More specifically, if we let $\delta(n)$ denote the number of divisors of the integer n , then—see Fig. 4—

$$\mathcal{H}(x, y) = \sum_{k=1}^{xy-1} \delta(k) + \begin{array}{l} \text{the position of } \langle x, y \rangle \text{ among} \\ \text{2-part factorizations of } xy, \\ \text{in reverse lexicographic order} \end{array} \quad (3.4)$$

1	3	5	8	10	14	16	...
2	7	13	19	26	34	40	...
4	12	22	33	44	56	69	...
6	18	32	48	64	81	99	...
9	25	43	63	86	108	130	...
11	31	55	80	107	136	165	...
15	39	68	98	129	164	200	...
17	47	79	116	154	193	235	...
⋮	⋮	⋮	⋮	⋮	⋮	⋮	⋮

Figure 4: The hyperbolic PF \mathcal{H} . The shell $xy = 6$ is highlighted.

The optimality of the PF \mathcal{H} in compactness (up to constant factors) is easily seen via the following argument. The set of arrays that have n or fewer positions are those of aspect ratios $a_i \times b_i$, where $a_i b_i \leq n$. As one sees from Fig. 5 (generalized to arbitrary n), the union of the positions of all these arrays is the set of integer lattice points under the hyperbola $xy = n$. Easily, this set of points has cardinality

$\Theta(n \log n)$. Since every array contains position $\langle 1, 1 \rangle$, it follows that, for every n , some array containing n or fewer positions is spread over $\Omega(n \log n)$ addresses.




Figure 5: *The aggregate set of positions of arrays having 16 or fewer position.*

Aside. The work described in this section aimed at giving one a broad range of ways of accessing one’s arrays/tables: by position, by row/column, by block (at varying computational costs). If one is interested in accessing an extendible array/table only by position, then one might be well served by the hashing schemes studied in [14]. In that source, we crafted hashing schemes with the following resource consumption. When one’s array/table has at most n positions, then, no matter what the array/table’s aspect ratio, the hashing scheme will employ fewer than $2n$ memory locations and will allow one to access any position of the array/table in expected time $O(1)$ and worst-case time $O(\log \log n)$.

4. Pairing Functions and Web-Computing

Evolving technology has given rise to a new modality of cooperative computing, which we call *Web-Based Computing* (WBC, for short). WBC proceeds roughly as follows. “Volunteers” register with a WBC website. After having registered, each volunteer visits the website from time to time to receive a task to compute. Some time after completing a task, the volunteer returns the results from that task and receives a new task. And the cycle continues.

As typically implemented, WBC is vulnerable to malicious, or careless, volunteers returning false results. When the WBC computations relate to sensitive matters such as security [15] or clinical drug testing [6, 10], such false results could have dire consequences. In [13], I have proposed a computationally lightweight scheme for keeping track of which volunteer computed which task(s), thereby enabling the head of the WBC project to ban frequently errant volunteers from continued partic-

ipation in the project. Note that my work addresses concerns about *accountability*, *not security*.

The scheme I proposed builds on the strategy of assigning positive-integer indices to (a) the set of all tasks, (b) all volunteers, (c) the set of tasks reserved for each volunteer v , and using a PF \mathcal{T} (which I call a *task-allocation function* in this context) to link volunteers with their assigned tasks. In other words, the t th task that volunteer v receives to compute is task $\mathcal{T}(v, t)$. Since the potential practicality of such a scheme demands that the functions \mathcal{T} , \mathcal{T}^{-1} , and $\mathcal{S}(v, t) \stackrel{\text{def}}{=} \mathcal{T}(v, t+1) - \mathcal{T}(v, t)$ all be easily computed, the primary focus in [13] is on PFs that are *additive* (APFs, for short): an APF assigns each volunteer v a *base task-index* B_v and a *stride* S_v ; it then uses the formula

$$\mathcal{T}(v, t) = B_v + (t - 1)S_v$$

to determine the workload task-index of the t th task assigned to volunteer v . From a system perspective, APFs have the benefit that a volunteer’s stride need be computed only when s/he registers at the website and can be stored for subsequent appearances.

A task-allocation scheme based entirely on APFs allows new volunteers to arrive dynamically but not to depart. If a volunteer departs, his/her tasks will never be computed—unless a new volunteer arrives to take their places and compute their tasks. Such reassignment would demand added mechanisms to retain accountability. The complete scheme described in [13] has a “front end” which allows volunteers to arrive *and depart* dynamically; it also ensures that faster volunteers are always assigned smaller indices. The interested reader is referred to that paper for details beyond this section’s focus on APFs.

Given the proposed use of APFs to assign indices to volunteers, one can argue that the management of the memory where tasks reside is simplified if one devises APFs whose strides S_v grow slowly as a function of v . Such APFs are “compact,” in the sense of (3.1). This observation sets the agenda for [13] and for the remainder of this section. Section 4.1 presents a methodology for designing easily computed APFs; Section 4.2 presents a sequence of APFs that suggest a tradeoff between the ease of computing an APF and the rate of growth of the APF’s strides.

We henceforth abstract the preceding discussion from the WBC scenario by replacing “volunteer” by “row” and “base task-index” by “base row-entry.” We also revert to our generic uses of x and y , instead of v and t .

4.1. A Methodology for Designing Additive PFs

Easily, any APF must have infinitely many distinct strides; i.e., S_x , viewed as a function of x , must have infinite range. Despite this, that there do exist easily computed APFs. One strategy for designing such APFs builds on the following well-known property of the set \mathcal{O} of positive odd integers.

Lemma 4.1 ([9]) *For any positive integer c , every odd integer can be written in precisely one of the 2^{c-1} forms:*

$$2^c n + 1, 2^c n + 3, 2^c n + 5, \dots, 2^c n + (2^c - 1),$$

for some nonnegative integer n .

One builds on Lemma 4.1 to construct APFs as follows.

Procedure APF-Constructor(\mathcal{T})

/*Construct an APF \mathcal{T}^* /*

Step 1. Partition the set of row-indices into *groups* whose sizes are powers of 2 (with any desired mix of equal-size and distinct-size groups). Order the groups linearly in some (arbitrary) way.

/*One can now talk unambiguously about group 0 (whose members share *group-index* $g = 0$), group 1 (whose members share group-index $g = 1$), and so on.*/*

Step 2. Assign each group a distinct copy of the set \mathcal{O} , as well as a *copy-index* $\kappa(g)$ expressed as a function of the group-index g .

Step 3. Allocate group g 's copy of \mathcal{O} to its members via the ($c = \kappa(g)$) instance of Lemma 4.1, using the multiplier 2^g as a *signature* to distinguish group g 's copy of the set \mathcal{O} from all other groups' copies.

Procedure APF-Constructor can be viewed as specializing the general scheme for constructing APFs in [16]. The specialization allows us to specify the APF in a computationally friendly way.

An explicit expression for \mathcal{T} . If we denote the $2^{\kappa(g)}$ rows of group g as $x_{g,1}, x_{g,2}, \dots, x_{g,2^{\kappa(g)}}$, then for all $i \in \{1, 2, \dots, 2^{\kappa(g)}\}$,

$$\mathcal{T}(x_{g,i}, y) \stackrel{\text{def}}{=} 2^g \left[2^{1+\kappa(g)}(y-1) + (2x_{g,i} + 1 \bmod 2^{1+\kappa(g)}) \right] \quad (4.1)$$

Theorem 4.2 Any function $\mathcal{T} : \mathbb{N} \times \mathbb{N} \leftrightarrow \mathbb{N}$ that is designed via Procedure APF-Constructor, hence is of the form (4.1), is a valid APF whose base row-entries and strides satisfy

$$B_x \leq S_x = \mathcal{T}(x, y+1) - \mathcal{T}(x, y) = 2^{1+g+\kappa(g)}. \quad (4.2)$$

Proof. We sketch the proof. (1) Any such \mathcal{T} maps $\mathbb{N} \times \mathbb{N}$ onto \mathbb{N} , because every positive integer equals *some* power of 2 times *some* odd integer. \mathcal{T} is *one-to-one* because it has a functional inverse \mathcal{T}^{-1} . To wit, the trailing 0's of each image integer $k = \mathcal{T}(x, y)$ identify x 's group g , hence the operative instance $\kappa(g)$ of Lemma 4.1. Then:

1. We compute

$$x = \frac{1}{2} \left[(2^{-g} k \bmod 2^{1+\kappa(g)}) - 1 \right],$$

which is an integer because the division by 2^g produces an odd number.

2. This leaves us with a linear expression of the form $ay + b$, from which we easily compute y .

Finally, we read the relations (4.2) directly from (4.1). ■

In order to implement Procedure APF-Constructor completely, one must express both the group-indices g and their associated copy-indices $\kappa(g)$ as functions of x . This is accomplished by noting that all x whose indices lie in the range

$$2^{\kappa(0)} + 2^{\kappa(1)} + \dots + 2^{\kappa(g-1)} + 1 \leq x \leq 2^{\kappa(0)} + 2^{\kappa(1)} + \dots + 2^{\kappa(g-1)} + 2^{\kappa(g)} \quad (4.3)$$

share group-index g and copy-index $\kappa(g)$. Translating the range (4.3) into an efficiently computed expression of the form $g = f(x)$ may be a simple or a challenging enterprise, depending on the functional form of $\kappa(g)$ that results from the grouping of row-indices.

4.2. A Sampler of Explicit APFs

Theorem 4.2 assures us that Procedure APF-Constructor produces a valid APF no matter how the copy-index $\kappa(g)$ grows as a function of the group-index g . However, the ease of computing the resulting APF, and its compactness, depend crucially on this growth rate. We now illustrate how one can use this growth rate as part of the design process, in order to stress either the ease of computing an APF or its compactness.

4.2.1. APFs that stress computation ease

We first implement Procedure APF-Constructor with *equal-size groups*, i.e., with $\kappa(g) = \text{constant}$. For each $c \in \mathbb{N}$, let $\mathcal{T}^{<c>}$ be the APF produced by the Procedure with $\kappa^{<c>}(g) \equiv c - 1$. One computes easily that

$$\mathcal{T}^{<c>}(x, y) \stackrel{\text{def}}{=} 2^{\lfloor (x-1)/2^{c-1} \rfloor} [2^c(y-1) + (2x-1 \bmod 2^c)].$$

Proposition 4.1 *Each $\mathcal{T}^{<c>}$ is a valid APF whose base row-entries and strides are given by*

$$B_x^{<c>} \leq S_x^{<c>} = 2^{\lfloor (x-1)/2^{c-1} \rfloor + c}. \quad (4.4)$$

Each $\mathcal{T}^{<c>}$ is easy to compute but has base row-entries and strides that grow *exponentially* with row-indices. Increased values of c (= larger fixed group sizes) decrease the base of the growth exponential, at the expense of modest increase in computational complexity. Computing a few sample values illustrates how a larger value of c penalizes a few low-index rows but gives all others significantly smaller base row-entries and strides; cf. the top half of Fig. 6.

4.2.2. APFs that balance computation ease and compactness

The functional form of the exponent of 2 in (4.4) suggests that one can craft an APF whose base row-entries and strides grow subexponentially by allowing the parameter c to grow with x , in a way that (roughly) balances $x/2^c$ against c . This strategy leads us to consider the copy-index $\kappa^\#(g) = g$. When we implement Procedure APF-Constructor with copy-index $\kappa^\#$, we arrive at an APF $\mathcal{T}^\#$ that is rather

x	g	$\mathcal{T}^{<1>}(x, y)$					
14	13	8192	24576	40960	57344	73728	...
15	14	16384	49152	81920	114688	147456	...
x	g	$\mathcal{T}^{<3>}(x, y)$					
14	3	24	88	152	216	280	...
15	3	40	104	168	232	296	...
		\vdots	\vdots	\vdots	\vdots	\vdots	\vdots
28	6	448	960	1472	1984	2496	...
29	7	128	1152	2176	3200	4224	...
x	g	$\mathcal{T}^{\#}(x, y)$					
28	4	400	912	1424	1936	2448	...
29	4	432	944	1456	1968	2480	...
x	g	$\mathcal{T}^{\star}(x, y)$					
28	3	328	840	1352	1864	2376	...
29	3	344	856	1368	1880	2392	...
		\vdots	\vdots	\vdots	\vdots	\vdots	\ddots

Figure 6: *Sample values by several APFs.*

easy to compute and whose base row-entries and strides grow only *quadratically* with row-indices. To wit ...

The copy-index $\kappa^{\#}(g) = g$ aggregates row-indices into groups of exponentially growing sizes: each group g comprises row-indices $2^g, 2^g + 1, \dots, 2^{g+1} - 1$. By (4.3), then, one computes easily that^a

$$\kappa^{\#}(g) = g = \lfloor \log x \rfloor. \quad (4.5)$$

Instantiating (4.5) in the definitional scheme (4.1), we find that

$$\mathcal{T}^{\#}(x, y) = 2^{\lfloor \log x \rfloor} \left(2^{1+\lfloor \log x \rfloor} (y - 1) + (2x + 1 \bmod 2^{1+\lfloor \log x \rfloor}) \right) \quad (4.6)$$

Proposition 4.2 *The function $\mathcal{T}^{\#}$ specified in (4.6) is a valid APF whose base row-entries and strides (as functions of x) are given by*

$$B_x^{\#} < \mathcal{S}_x^{\#} = 2^{1+2\lfloor \log x \rfloor} \leq 2x^2,$$

hence, grow quadratically with x .

Comparing $\mathcal{T}^{\#}$ and the $\mathcal{T}^{<c>}$. For sufficiently large x , the (exponentially growing) strides of any of the APFs $\mathcal{T}^{<c>}$ will be dramatically larger than the (quadratically growing) strides of the APF $\mathcal{T}^{\#}$. However, it takes a while for $\mathcal{T}^{\#}$'s superiority to manifest itself; for instance,

- it is not until $x = 5$ that $\mathcal{T}^{<1>}$'s strides are always at least as large as $\mathcal{T}^{\#}$'s;

^aThroughout, all logarithms have base 2.

- the corresponding number for $\mathcal{T}^{<2>}$ is $x = 11$;
- the corresponding number for $\mathcal{T}^{<3>}$ is $x = 25$.

4.2.3. APFs that stress compactness

By choosing a copy-index $\kappa(g)$ that grows superlinearly with g , one can craft APFs whose base row-entries and strides grow subquadratically, thereby beating the compactness of $\mathcal{T}^\#$. But one must choose $\kappa(g)$'s growth rate judiciously, since faster growth need not enhance compactness.

Achieving subquadratic growth. Many copy-index growth rates yield APFs with subquadratic compactness. However, all of the APFs we know of that achieve this goal are rather difficult to compute and actually achieve the goal only asymptotically, hence are more likely of academic than of practical interest.

Consider, for each $k \in \mathbb{N}$, the APF $\mathcal{T}^{[k]}$ specified by the copy-index $\kappa^{[k]}(g) = g^k$. By (4.3), the row-indices x belonging to group g now lie in the range

$$1 + 2 + 2^{2^k} + \cdots + 2^{(g-1)^k} < x \leq 1 + 2 + 2^{2^k} + \cdots + 2^{g^k},$$

so that $g = (1 + o(1))\lceil(\log x)^{1/k}\rceil$. We actually use the simplified, albeit slightly inaccurate, expression $g = \lceil(\log x)^{1/k}\rceil$ in our asymptotic analyses of the $\mathcal{T}^{[k]}$, since the $o(1)$ -quantity decreases very rapidly with growing x . Although closed-form expressions for $\mathcal{T}^{[k]}$ in terms of x have eluded us, we can verify that each $\mathcal{T}^{[k]}$ does indeed enjoy subquadratic stride growth.

Proposition 4.3 *Each function $\mathcal{T}^{[k]}$ produced by Procedure APF-Constructor from the copy-index $\kappa^{[k]}(g) = g^k$ is a valid APF whose base row-entries and strides (as functions of x) are given by*

$$B_x^{[k]} \leq \mathcal{S}_x^{[k]} = 2^{O((\log x)^{1/k} + \log x)} = x 2^{O((\log x)^{1/k})} \quad (4.7)$$

hence, grow subquadratically with x .

We illustrate a close relative of $\mathcal{T}^{[2]}$ which exhibits its subquadratic compactness at much smaller values of x than $\mathcal{T}^{[2]}$ does, namely, the APF \mathcal{T}^* that Procedure APF-Constructor produces from the copy-index

$$\kappa^*(g) = \left\lceil \frac{1}{2} g^2 \right\rceil. \quad (4.8)$$

Mimicking the development with $\kappa^{[k]}$, we see that now $g = (1 + o(1))\lceil\sqrt{2 \log x}\rceil + 1$, which we simplify for analysis to the slightly inaccurate expression

$$g = \left\lceil \sqrt{2 \log x} \right\rceil + 1.$$

We can easily compute \mathcal{T}^* from (4.8), in the presence of (4.1, 4.3).

Proposition 4.4 *The base row-entries and strides of the APF \mathcal{T}^* satisfy*

$$B_x^* \leq \mathcal{S}_x^* = 2^{1+g+\kappa^*(g)} \approx 8x4^{\sqrt{2 \log x}}.$$

Comparing \mathcal{T}^* and $\mathcal{T}^\#$. Any function that grows quadratically with x will eventually produce significantly larger values than a function that grows only as $x4\sqrt{2\log x}$. Therefore, \mathcal{T}^* 's strides will eventually be dramatically smaller than $\mathcal{T}^\#$'s. Fig. 6 indicates that this difference takes effect at about the same point as the exponential vs. quadratic one noted earlier, albeit at the cost of greater computational complexity.

The danger of excessively fast growing κ . If $\kappa(g)$ grows too fast with g , then the base row-entries and strides of the resulting APF grow *superquadratically* with the row-indices x , thereby confuting our goal of beating quadratic growth. We exemplify this fact by supplying Procedure APF-Constructor with the copy-index $\kappa(g) = 2^g$; the reader can readily supply other examples. By (4.3), we see that in this case, $g = \lfloor \log \log x \rfloor + O(1)$. Therefore, whenever x is the smallest row-index with a given group-index g (of course, infinitely many such x exist) we have

$$x = 2^{\kappa(0)} + 2^{\kappa(1)} + \dots + 2^{\kappa(g-1)} + 1 \approx \sqrt{2^{\kappa(g)}},$$

while the stride associated with x is (cf. (4.2))

$$S_x = 2^{1+g+\kappa(g)} > 2^{\kappa(g)} \kappa(g) \approx x^2 \log x.$$

We do not yet know the growth rate at which faster growing $\kappa(g)$ starts hurting compactness. Finding this rate is an attractive research problem.

Acknowledgements

The research described in Sections 2 and 3 was done while the author was with the IBM Watson Research Center. The research described in Section 4, and the preparation of this paper were supported in part by NSF Grant CCR-00-73401.

References

1. G. Cantor (1878): Ein Beitrag zur Begründung der transfiniten Mengenlehre. *J. Reine Angew. Math.* 84, 242–258.
2. A.L. Cauchy (1821): *Cours d'analyse de l'École Royale Polytechnique, 1ère partie: Analyse algébrique*. l'Imprimerie Royale, Paris. Reprinted: Wissenschaftliche Buchgesellschaft, Darmstadt, 1968.
3. M. Davis (1958): *Computability and Unsolvability*. McGraw-Hill, N.Y.
4. R. Fueter and G. Pólya (1923): Rationale Abzählung der Gitterpunkte. *Vierteljschr. Naturforsch. Ges. Zürich* 58, 380–386.
5. K. Gödel (1931): Über formal unentscheidbare Sätze der Principia Mathematica und verwandter Systeme, I. *Monatshefte für Mathematik und Physik* 38, 173–198.
6. *The Intel Philanthropic Peer-to-Peer program*. (www.intel.com/cure).
7. J.S. Lew and A.L. Rosenberg (1978): Polynomial indexing of integer lattices, I. *J. Number Th.* 10, 192–214.
8. J.S. Lew and A.L. Rosenberg (1978): Polynomial indexing of integer lattices, II. *J. Number Th.* 10, 215–243.

9. I. Niven and H.S. Zuckerman (1980): *An Introduction to the Theory of Numbers*. (4th Ed.) J. Wiley & Sons, New York.
10. *The Olson Laboratory Fight AIDS@Home project*. (www.fightaidsathome.org).
11. A.L. Rosenberg (1974): Allocating storage for extendible arrays. *J. ACM* 21, 652–670.
12. A.L. Rosenberg (1975): Managing storage for extendible arrays. *SIAM J. Comput.* 4, 287–306.
13. A.L. Rosenberg (2003): Accountable Web-computing. *IEEE Trans. Parallel and Distr. Sys.* 14, to appear.
14. A.L. Rosenberg and L.J. Stockmeyer (1977): Hashing schemes for extendible arrays. *J. ACM* 24, 199–221.
15. *The RSA Factoring by Web Project*. (<http://www.npac.syr.edu/factoring>) (with Foreword by A. Lenstra). Northeast Parallel Architecture Center.
16. L.J. Stockmeyer (1973): Extendible array realizations with additive traversal. IBM Research Report RC-4578.
17. A.M. Turing (1936): On computable numbers, with an application to the Entscheidungsproblem. *Proc. London Math. Soc.* (ser. 2, vol. 42) 230–265; Correction *ibid.* (vol. 43) 544–546.